

Assignment 3: Image Processing

(1)

I uploaded an image from my computer to my Jupyter Notebook home page. Here is the image in a markdown cell:



(2)

Now let's write a python code to import the image and export some basic information.

```
In [1]: from PIL import Image  
        bs = Image.open('baby_snek.jpg')  
        print(bs.size)  
        print(bs.mode)
```

```
(838, 559)  
RGB
```


The image is opened and we see that it is a 838 pixels wide by 559 pixels high and is in RGB mode.

(3)

In order to get a better understanding of edge detection let's try the Image Filter FIND_EDGES

```
In [2]:  ▶ from PIL import ImageFilter
        bs2 = bs.filter(ImageFilter.FIND_EDGES)
        bs2.show()
```

Here is the resulting image:

**(4)**

In order to reduce noise we will apply a blurring filter.

```
In [3]:  ▶ bs3 = bs.filter(ImageFilter.GaussianBlur(radius=2))
        bs3.show()
        bs4 = bs.filter(ImageFilter.GaussianBlur(radius=5))
        bs4.show()
        bs5 = bs.filter(ImageFilter.GaussianBlur(radius=10))
        bs5.show()
```

Here is a Gaussian blur with radius 2 pixels



And the same filter with radius 5 pixels



And with radius 10 pixels



```
In [4]: ▶ bs6 =bs5.filter(ImageFilter.FIND_EDGES)  
        bs6.show()
```

If the image is too blurred FIND_EDGES doesn't detect anything



Let's try to make our own Gaussian blurring filter from scratch. Pixels at the edge of the image we will call boundary pixels. They don't have enough surrounding pixels to run some calculations. For the blurring step boundary pixels keep their original values.

A Gaussian filter of the size $(2k + 1)^2$ can be defined by the equation:

$$H_{ij} = \frac{1}{2\pi\sigma^2} e^{\left(-\frac{(i-(k+1))^2 + (j-(k+1))^2}{2\sigma^2}\right)}$$

Where $1 \leq i, j \leq (2k + 1)$

```
In [5]: import numpy as np

def GausKer(size,sig):
    k = int(size) // 2
    x,y = np.mgrid[-k:k+1, -k:k+1]
    gk = (1/(2*np.pi*sig**2))*(np.exp(-((x**2 + y**2) / (2*sig**2))))
    return gk

GausKer(5,10).flatten()
```

```
Out[5]: array([0.00152914, 0.00155225, 0.00156003, 0.00155225, 0.00152914,
               0.00155225, 0.00157571, 0.00158361, 0.00157571, 0.00155225,
               0.00156003, 0.00158361, 0.00159155, 0.00158361, 0.00156003,
               0.00155225, 0.00157571, 0.00158361, 0.00157571, 0.00155225,
               0.00152914, 0.00155225, 0.00156003, 0.00155225, 0.00152914])
```

```
In [6]: ▶ def BlurImg(size,sig,img):
        ker=GausKer(size,sig).flatten()
        img2 = img.filter(ImageFilter.Kernel((size,size), ker))
        return img2

        BlurImg(5,10,bs)
```

Out[6]:



(5)

Now we will make our own edge detection filter. We will use the first derivative of the Gaussian.

In []: ▶

```
In [7]: ▶ bsm= bs.convert('L')
        bsbw= bs.convert('1')
        bsm.show()
        bsbw.show()
```

```
In [8]: ▶ pixels = list(bsm.getdata())
        w,h = bsm.size
        pixels = [pixels[i*w:(i+1)*w]for i in range(h)]
        #print(pixels)
```

```
In [9]: ▶ pixels = np.asarray(pixels)
print(pixels.shape)
```


```
(559, 838)
```

```
In [10]: ▶ kern=[[0,-1,0],[-1,4,-1],[0,-1,0]]
```

```
In [11]: ▶ from scipy import ndimage
%matplotlib inline
from matplotlib import pyplot as plt

conv=ndimage.convolve(pixels,kern, mode='constant', cval=0.0)
img = Image.fromarray(conv)
img.show()
```

Here is our resulting image



Let's see if we can improve it by adding blurs with different widths

```
In [12]: ▶ bs2 = BlurImg(5,2,bs)
bs2m = bs2.convert('L')
pixels2 = list(bs2m.getdata())
w,h = bs2m.size
pixels2 = [pixels2[i*w:(i+1)*w] for i in range(h)]
pixels2 = np.asarray(pixels2)
conv2=ndimage.convolve(pixels2,kern, mode='constant', cval=0.0)
img2 = Image.fromarray(conv2)
img2.show()
```

Here is our resulting image for $\sigma = 2$, it is much harder to see and smaller details are smoothed out.





```
In [13]: ▶ def EdgeMe2D(img):  
    pixels2 = list(img.getdata())  
    w,h = img.size  
    pixels2 = [pixels2[i*w:(i+1)*w]for i in range(h)]  
    pixels2 = np.asarray(pixels2)  
    conv2=ndimage.convolve(pixels2,kern, mode='constant', cval=0.0)  
    img2 = Image.fromarray(conv2)  
    img2.show()  
    return img2  
  
def EdgeMe3D(img):  
    pixels2 = list(img.getdata())  
    w,h = img.size  
    pixels2 = [pixels2[i*w:(i+1)*w]for i in range(h)]  
    pixels2 = np.asarray(pixels2)  
    pixels2d= pixels2[:, :, 0]  
    conv2=ndimage.convolve(pixels2d,kern, mode='constant', cval=0.0)  
    img2 = Image.fromarray(conv2)  
    img2.show()  
    return img2
```

```
In [14]: ▶ from PIL import ImageEnhance  
def enhancer(img):  
    enhancer = ImageEnhance.Contrast(img)  
    enhanced_im = enhancer.enhance(4.0)  
    enhanced_im.show()  
    return enhanced_im
```

```
In [15]: EdgeMe3D(enhancer(bs2))
```

Out[15]:




If we enhance the contrast we can see that increasing the blur gives a nice outline of the snake but most of the detailed patterns are lost.



```
In [16]: ▶ bs100 = BlurImg(5,100,bs)  
EdgeMe3D(enhancer(bs100))
```

Out[16]:



We see little difference when increasing σ from 2 to 100.




(6)

Now let's play with different parameters to optimise the system.

We'll start by trying different colors.


```
In [17]: ▶ bsbw= bs.convert('1')  
          bsm= bs.convert('L')  
  
          %matplotlib inline  
          from matplotlib import pyplot as plt  
  
          plt.imshow(bs)
```

Out[17]: <matplotlib.image.AxesImage at 0x226a7ea7ec8>




```
In [18]: ▶ plt.imshow(EdgeMe3D(bs))
```

Out[18]: <matplotlib.image.AxesImage at 0x226a7ec5fc8>




```
In [19]: ▶ plt.imshow(bsbw)
```

```
Out[19]: <matplotlib.image.AxesImage at 0x226a5d23648>
```




```
In [20]: ▶ plt.imshow(EdgeMe2D(bsbw))
```

```
Out[20]: <matplotlib.image.AxesImage at 0x226a5d92848>
```




```
In [21]: plt.imshow(EdgeMe2D(enhancer(bs).convert('1')))
```


```
Out[21]: <matplotlib.image.AxesImage at 0x226a5e00808>
```



The black and white image remains faint even when enhanced


```
In [22]: plt.imshow(bsm)
```

```
Out[22]: <matplotlib.image.AxesImage at 0x226a63ce248>
```




```
In [23]: ▶ plt.imshow(EdgeMe2D(bsm))
```

```
Out[23]: <matplotlib.image.AxesImage at 0x226a642b948>
```



```
In [24]: ▶ plt.imshow(EdgeMe2D(enhancer(bs).convert('L')))
```

```
Out[24]: <matplotlib.image.AxesImage at 0x226a648ef88>
```




The clearest image is derived from the RGB mode image.

We can also try changing the stencil we used.

```
In [25]: ▶ def EdgeMe3DKern2(img):  
    pixels2 = list(img.getdata())  
    w,h = img.size  
    kern=[[-1,-1,-1],[-1,8,-1],[-1,-1,-1]]  
    pixels2 = [pixels2[i*w:(i+1)*w]for i in range(h)]  
    pixels2 = np.asarray(pixels2)  
    pixels2d= pixels2[:, :, 0]  
    conv2=ndimage.convolve(pixels2d,kern, mode='constant', cval=0.0)  
    img2 = Image.fromarray(conv2)  
    img2.show()  
    return img2  
  
plt.imshow(EdgeMe3DKern2(bs))
```

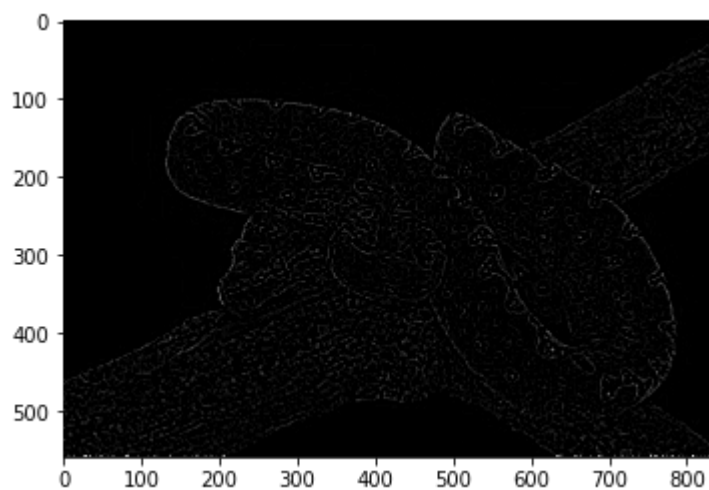
Out[25]: <matplotlib.image.AxesImage at 0x226a64c98c8>




lines show up brighter with this stencil, but there is also more noise

```
In [26]: ▶ plt.imshow(EdgeMe3DKern2(BlurImg(5,10,bs)))
```


```
Out[26]: <matplotlib.image.AxesImage at 0x226aa5d97c8>
```



Although this image is faint, if we look closely we can see that the edge detection is quite good. Small details in the pattern are detected, but there is less noise. Let's see if we can enhance the brightness of the image to make this more clear.



In order to avoid double detections we could add a non maximum suppression algorithm. This look at each pixel and checks to see if its neighbors $(i, j-1)$ and $(i, j+1)$ are brighter or dimmer. If a pixel is not a max, it is set to 0. This does not appear to be a major problem for our image. CAPTCHA stands for "completely automated public Turing test to tell computers and humans apart." The image is distorted in a way such that humans can decode it and type the answer, but computers cannot. For example this image is wavy but we can still easily read "smwm."



I think edge detection would be an important first step in decoding a CAPTCHA image, but the process would be very complex. We don't fully understand how the human brain can recognize letters as the same when they are distorted or in different fonts, so it would be difficult to program a computer to do so. I think it would have to involve a lot of machine learning. I think edge detection could have a lot of applications for visual artists. They could use it to identify and select different shapes in an image.

(7)

In this assignment we learned about basic image processing by convolving different kernels with a test image. By using a Gaussian kernel we created a blur filter. By using the Laplacian of the Gaussian we created an edge detection filter. This filter finds pixels that change intensity quickly

relative to their neighbors. The human eye detects this rapid change as an edge. Edge detection takes a complex image and can turn it into a simple line drawing. this has many applications for image processing.